

Introduction to Deep Learning

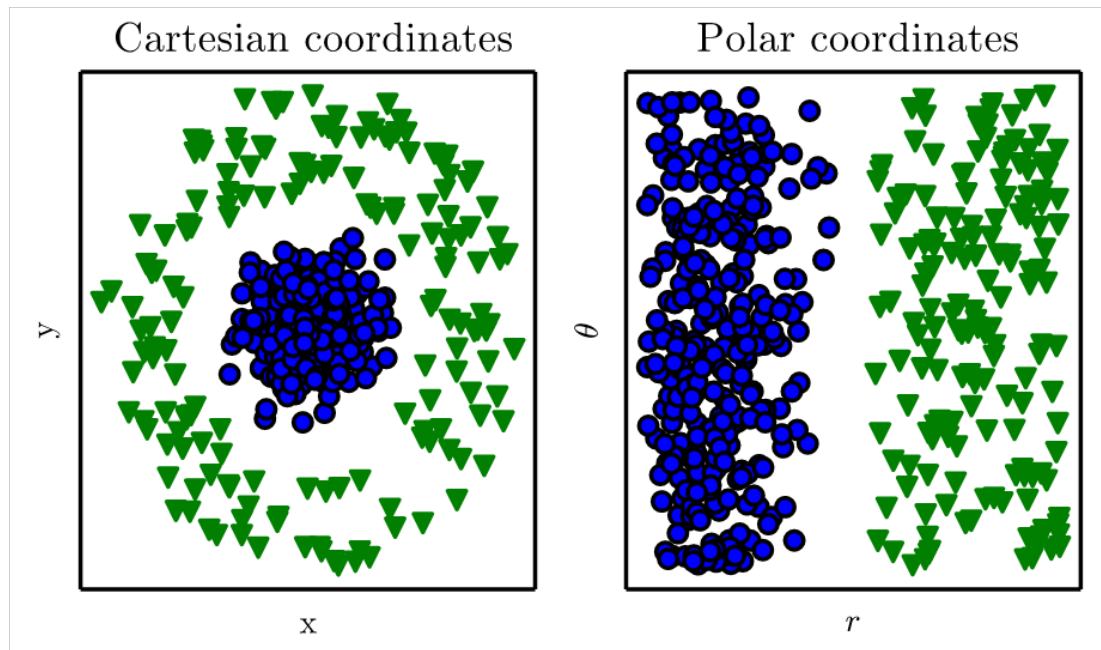
CMPT 733

Steven Bergner
sbergner@sfu.ca

Overview

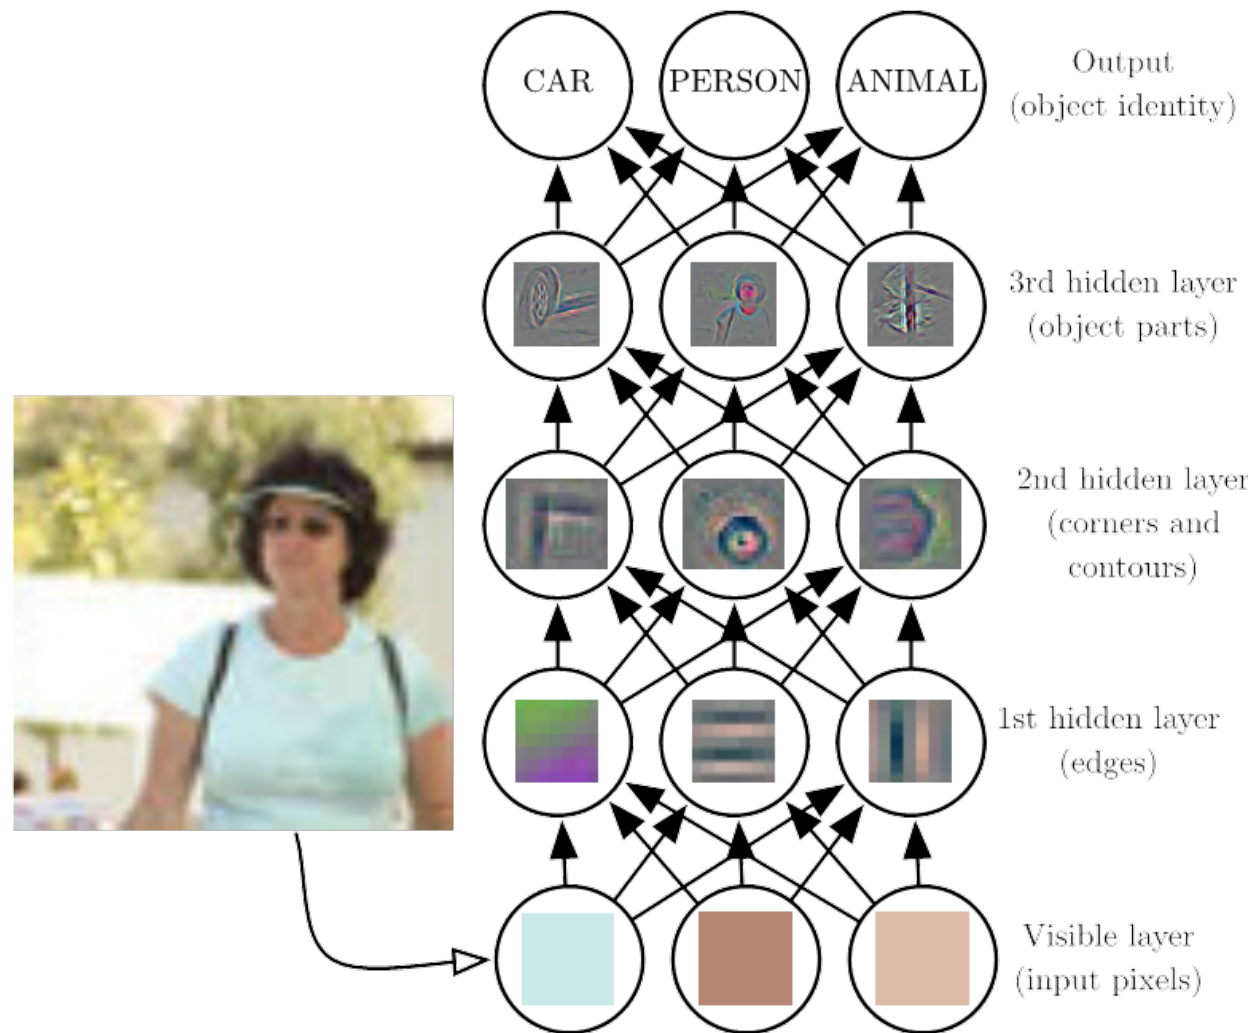
- Renaissance of artificial neural networks
 - Representation learning vs feature engineering
- Background
 - Linear Algebra, Optimization
 - Regularization
- Construction and training of layered learners
- Frameworks for deep learning

Representations matter

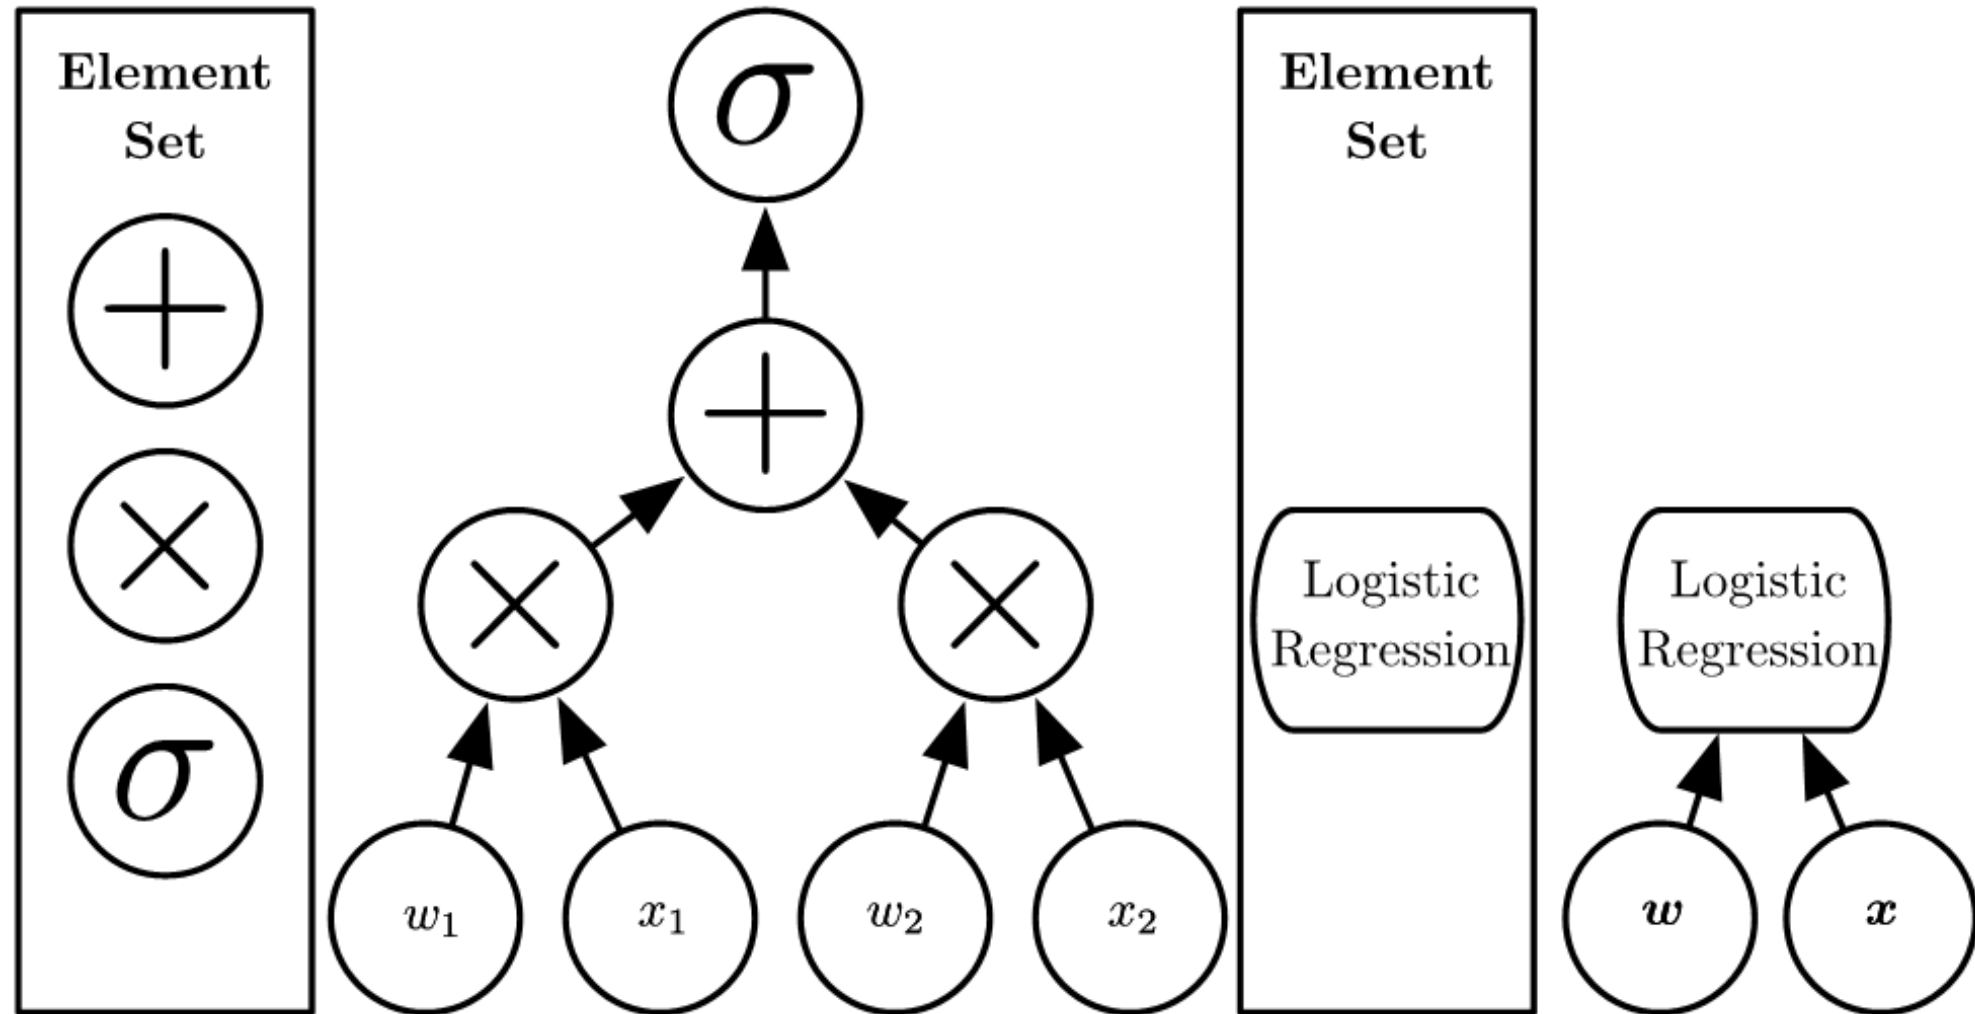


- Transform into the right representation
- Classify points simply by threshold on radius axis
- Single neuron with non-linearity can do this

Depth: layered composition

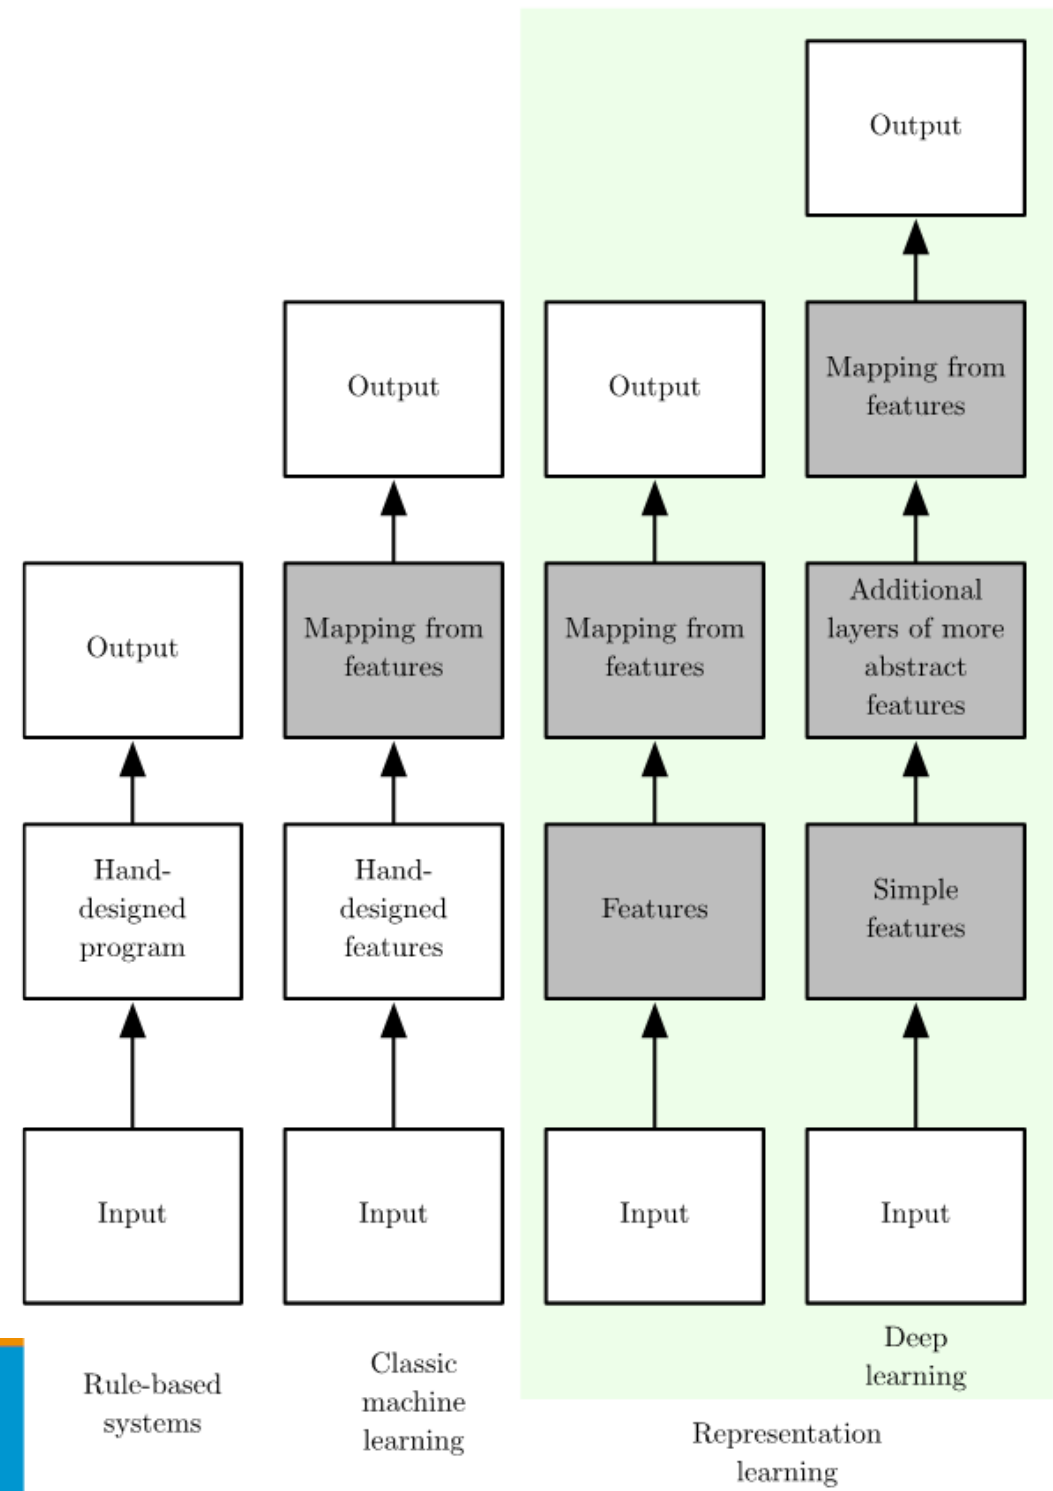


Computational graph




Components of learning

- Hand designed program
 - Input → Output
- Increasingly automated
 - Simple features
 - Abstract features
 - Mapping from features



Growing Dataset Size



MNIST dataset

8	9	0	1	2	3	4	7	8	9	0	1	2	3	4	5	6	7	8	6
4	2	6	4	7	5	5	4	7	8	9	2	9	3	9	3	8	2	0	5
0	1	0	4	2	6	5	3	5	3	8	0	0	3	4	1	5	3	0	8
3	0	6	2	7	1	1	8	1	7	1	3	8	9	7	6	7	4	1	6
7	5	1	7	1	9	8	0	6	9	4	9	9	3	7	1	9	2	2	5
3	7	8	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	0
1	2	3	4	5	6	7	8	9	8	1	0	5	5	1	9	0	4	1	9
3	8	4	7	7	8	5	0	6	5	5	3	3	3	9	8	1	4	0	6
1	0	0	6	2	1	1	3	2	8	8	7	8	4	6	0	2	0	3	6
8	7	1	5	9	9	3	2	4	9	4	6	5	3	2	5	5	9	4	1
6	5	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7
8	9	0	1	2	3	4	5	6	7	8	9	6	4	2	6	4	7	5	5
4	7	8	9	2	9	3	9	3	8	2	0	9	8	0	5	6	0	1	0
4	2	6	5	5	5	4	3	4	1	5	3	0	8	3	0	6	2	7	1
1	8	1	7	1	3	8	5	4	2	0	9	7	6	7	4	1	6	8	4
7	5	1	2	6	7	1	9	8	0	6	9	4	9	9	6	2	3	7	1
9	2	2	5	3	7	8	0	1	2	3	4	5	6	7	8	0	1	2	3
4	5	6	7	8	0	1	2	3	4	5	6	7	8	9	2	1	2	1	3
9	9	8	5	3	7	0	7	7	5	7	9	9	4	7	0	3	4	1	4
4	7	5	8	1	4	8	4	1	8	6	6	4	6	3	5	7	2	5	9


Basics

Linear Algebra and Optimization


Linear Algebra

- Tensor is an array of numbers
 - Multi-dim: 0d scalar, 1d vector, 2d matrix/image, 3d RGB image

- Matrix (dot) product $C = AB$ $C_{i,j} = \sum_k A_{i,k} B_{k,j}$



- Dot product of vectors A and B
 - ($m = p = 1$ in above notation, $n=2$)



Linear algebra: Norms

- L^p norm

$$||\mathbf{x}||_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$$


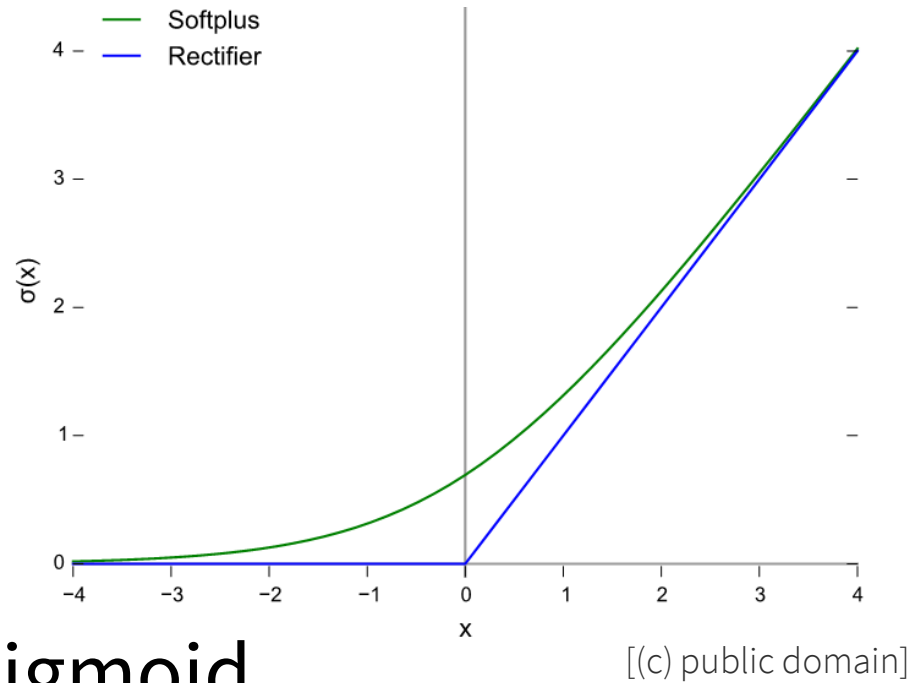
- Most popular norm: L2 norm, $p=2$
- L1 norm, $p=1$: $||\mathbf{x}||_1 = \sum_i |x_i|$.
- Max norm, infinite p : $||\mathbf{x}||_\infty = \max_i |x_i|$.

Nonlinearities


- ReLU

- Softplus


- Logistic Sigmoid



Approximate Optimization




Gradient descent




Critical points



Minimum




Maximum



Saddle point



Saddle point – 1 and 2 derivative vanish



Poor conditioning:
1 deriv large in one and small in another direction

Tensorflow Playground


- <http://playground.tensorflow.org/>
 - Try out simple network configurations
- <https://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>
 - Visualize linear and non-linear mappings

Regularization

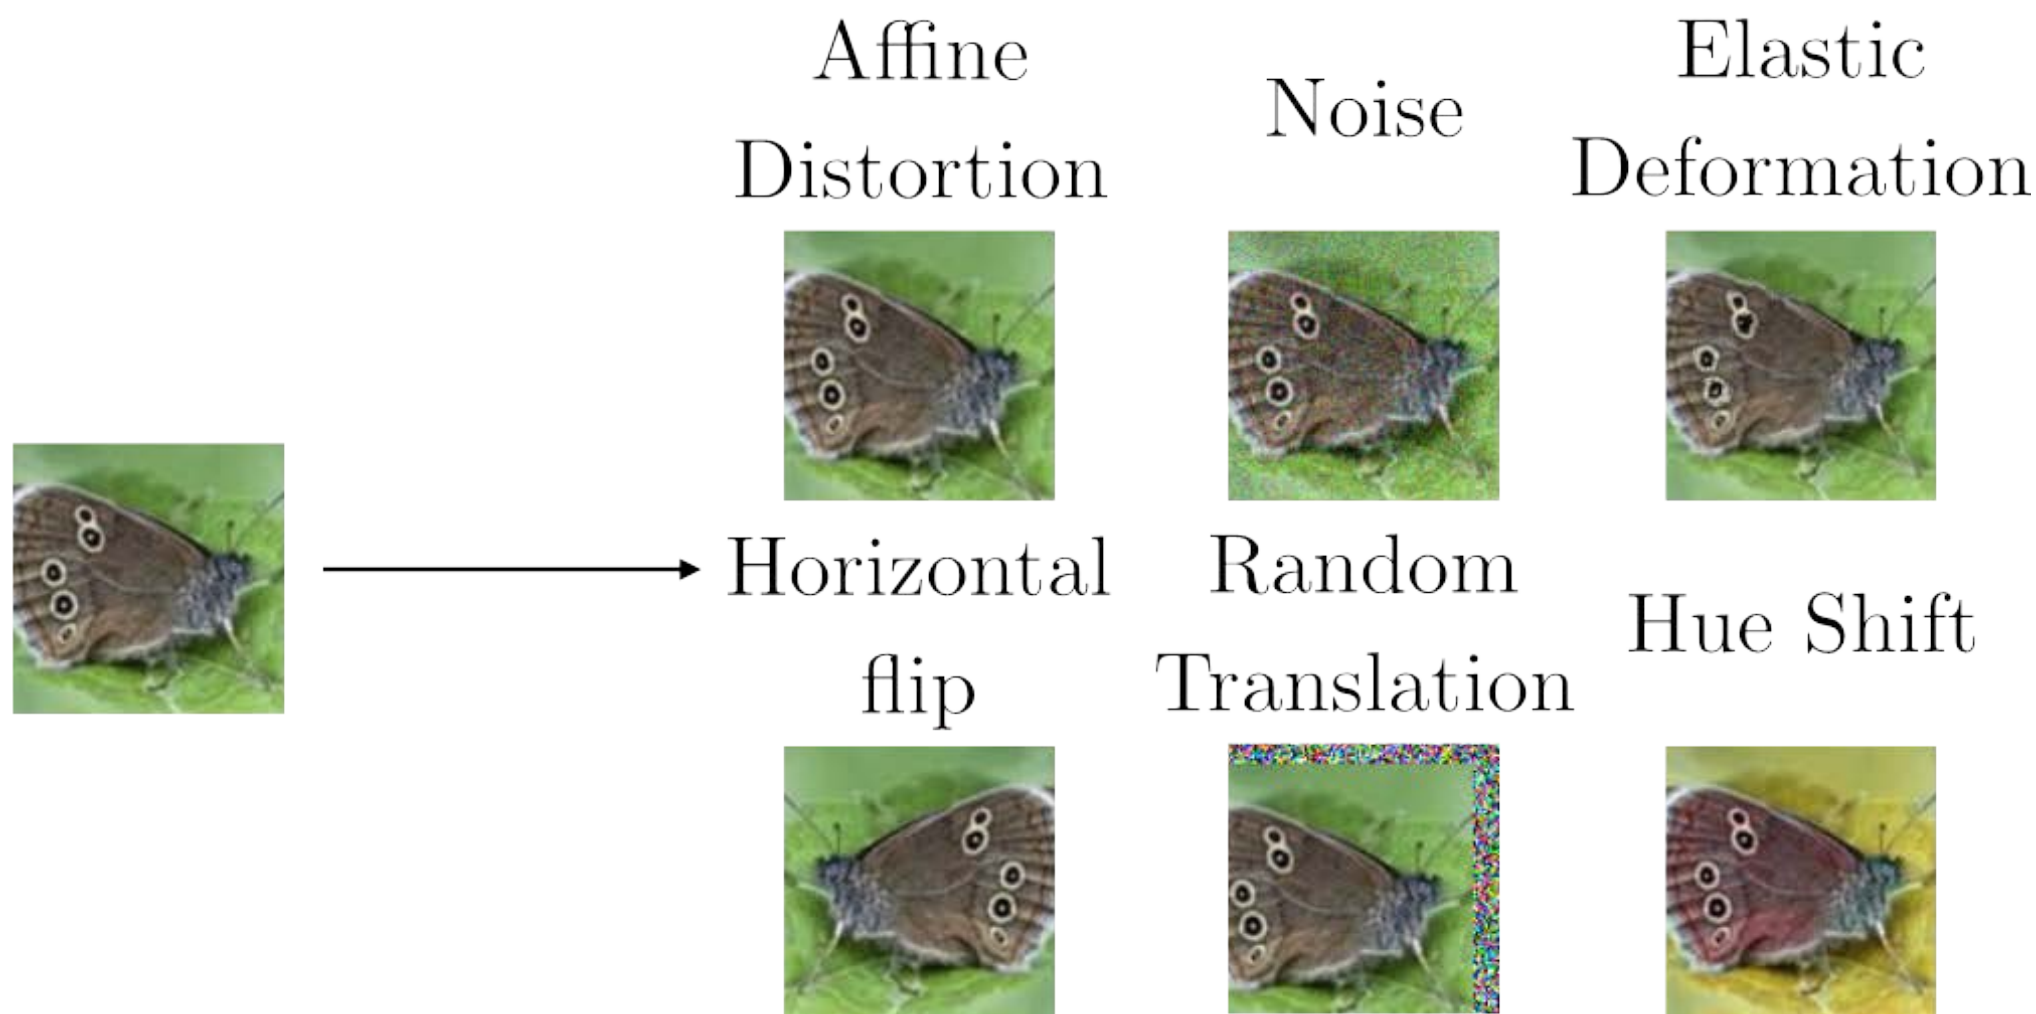
Reduced generalization error without impacting training error

Constrained optimization


- Squared L2 encourages small weights
- L1 encourages sparsity of model parameters (weights)



Dataset augmentation



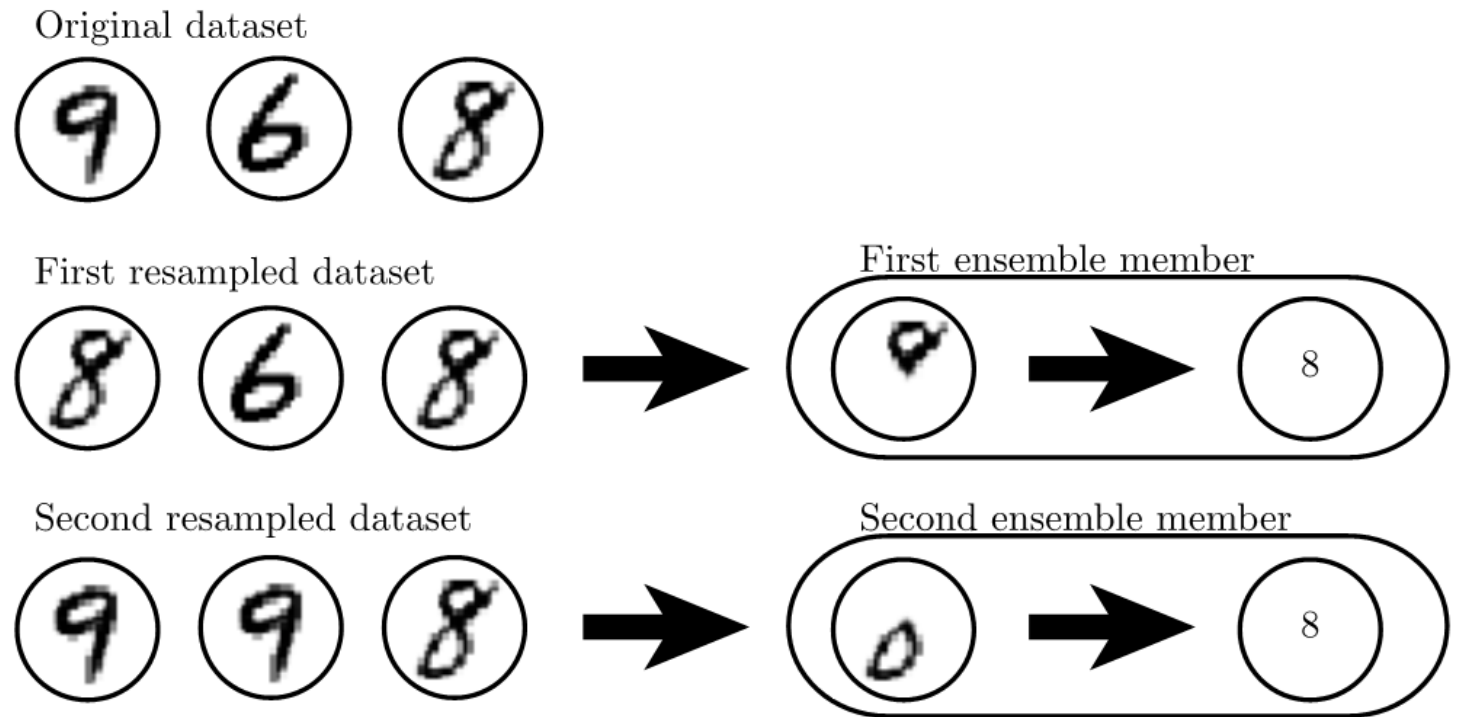
Learning curves



- Early stopping before validation error starts to increase



Bagging

- Average multiple models trained on subsets of the data
- First subset: learns top loop, Second subset: bottom loop




Dropout

- Random sample of connection weights is set to zero
- Train different network model each time
- Learn more robust, generalizable features



Multitask learning




- Shared parameters are trained with more data
- Improved generalization error due to increased statistical strength

Components of popular architectures


Convolution as edge detector



Input




Kernel




Output


Gabor wavelets (kernels)



Directional second
derivative




Second derivative
(curvature)




Local average, first
derivative

Gabor-like learned kernels




- Features extractors provided by pretrained networks

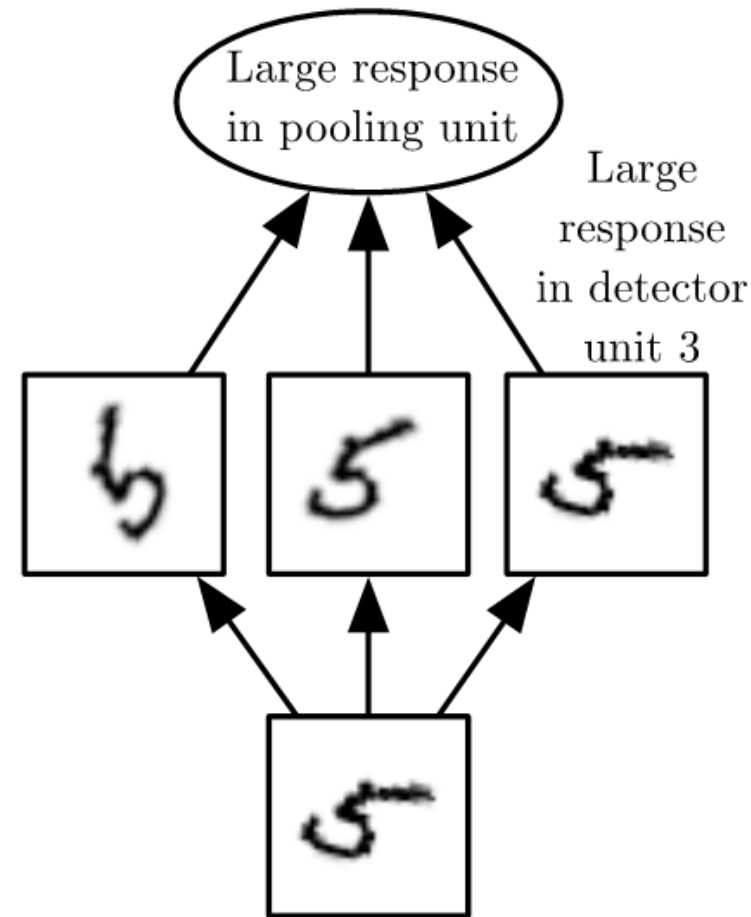
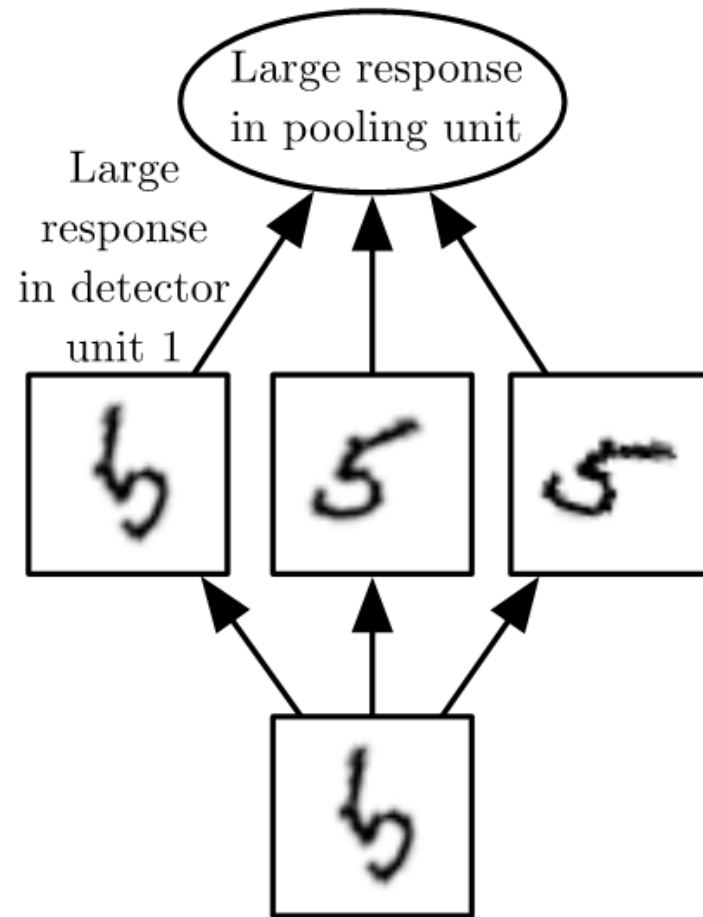
Max pooling translation invariance




- Take max of certain neighbourhood
- Often combined, followed by downsampling



Max pooling transform invariance



Types of connectivity



Local connection:
like convolution,
but no sharing

Choosing architecture family

- No structure → fully connected
- Spatial structure → convolutional
- Sequential structure → recurrent

Optimization Algorithm

- Lots of variants address choice of learning rate
- See [Visualization of Algorithms](#)
- AdaDelta and RMSprop often work well

Software for Deep Learning

Current Frameworks

- Tensorflow / Keras
- PyTorch
- DL4J
- Caffe (superseded by Caffe2, which is merged into PyTorch)
- [And many more](#)
- Most have CPU-only mode but much faster on NVIDIA GPU

Development strategy

- Identify needs: High accuracy or low accuracy?
- Choose metric
 - Accuracy (% of examples correct), Coverage (% examples processed)
 - Precision $TP/(TP+FP)$, Recall $TP/(TP+FN)$
 - Amount of error in case of regression
- Build end-to-end system
 - Start from baseline, e.g. initialize with pre-trained network
- Refine driven by data

Sources

- I. Goodfellow, Y. Bengio, A. Courville “Deep Learning” MIT Press 2016 [[link](#)]